
Introductory Seminar on Artificial Intelligence and Machine Learning

Emanuele Ledda, Cagliari Digital Lab 2024 - Day 1

About Us



Pattern Recognition
and Applications Lab



Emanuele Ledda

May 2024



Daniele Angioni

July 2024



Sara Concas

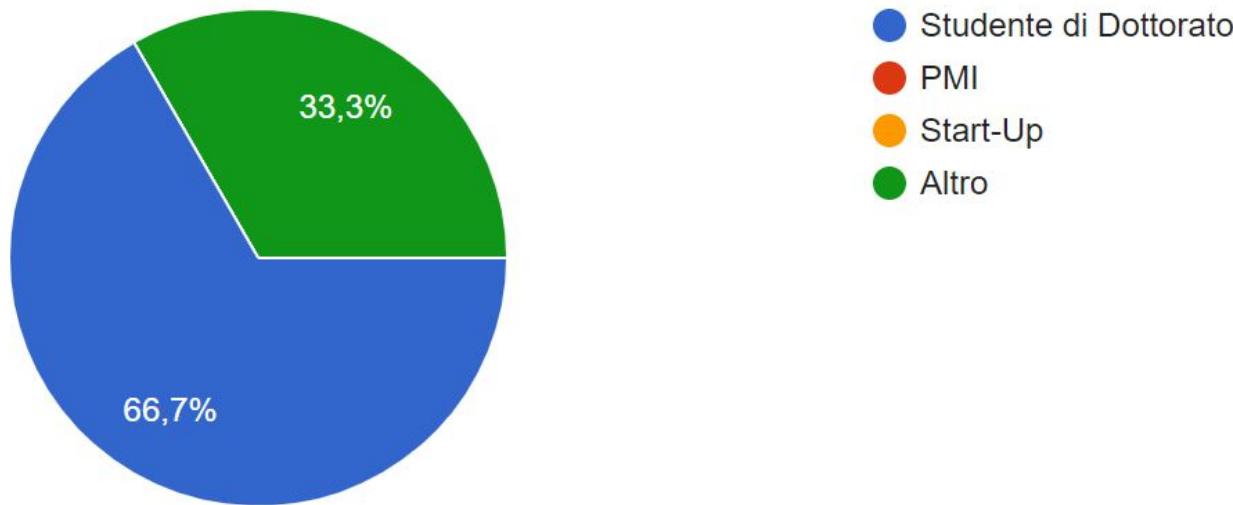
September 2024

Introductory Seminar on Artificial
Intelligence and Machine Learning

Introductory Seminar on PyTorch
for Deep Learning

Introductory Seminar on Computer
Vision

About You



Prerequisite

- Knowing at least one **programming language** (preferably **python**)
- Understanding the basic concepts of **Linear Algebra**
- Bring a **PC**



Seminar Objectives

- Become aware of the evolution of **Artificial Intelligence** over the years and how the central paradigm shifted from logical systems to statistical models until reaching modern **Machine Learning**.
- Learn how the most paramount **Shallow Learning** algorithms work, focusing on **Classification Algorithms**.
- Learn the functioning of **artificial neural networks**, from the perceptron to **deep neural networks**.
- Understand the fundamental concepts that led to the most famous modern **foundation models** (e.g., GPT 4.0, Sora)
- Understand the main critical issues behind **AI Ethics** and **Trustworthy AI** today, both from a technical and a regulatory perspective (**AI Act, European Guidelines for Trustworthy AI**).



Seminar Objectives - Technical Python Aspects



- Learning **numpy** essentials for managing tensor manipulation, essential for every ML task
- Basic data management skills with **pandas**
- General knowledge on using **SciKit-Learn** classifiers


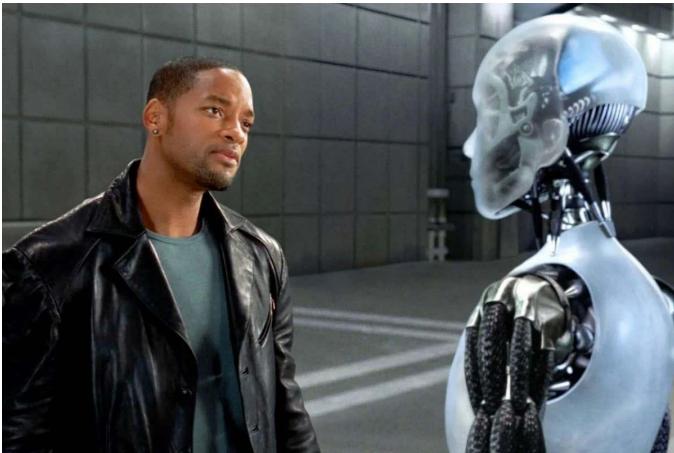


Introduction to Artificial Intelligence and Machine Learning




Artificial Intelligence

Imagining Artificial Intelligence



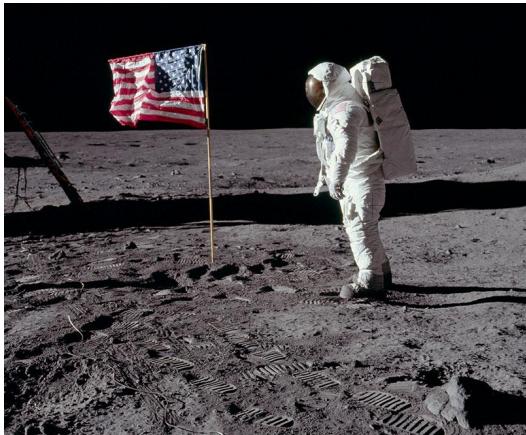
What is Intelligence?



- What should an intelligent agent **be**?
 - Should it be **similar to human beings**?
 - Should it be able to **play chess**?
 - Should it be able to express **artistic ideas**?



Human Intelligence can be this...



... But also this!

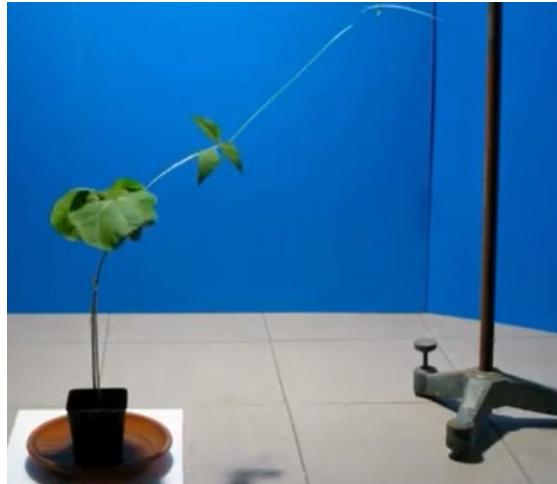


- But intelligence is **not** a human privilege...

Animal intelligence



Plant Intelligence



So, what is Intelligence?

- Intelligence does **NOT** mean to be similar to humans
- Intelligence is the ability of a system to **act appropriately** in an **uncertain environment**, where "appropriately" means "**maximizing the success probability.**"



Climbing = Gain more light



Dancing = Impress a partner



Cooking (originally) = Optimize the nutrients

Agent, Action, Environment and Sensory Input

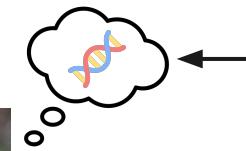
- act appropriately in an **uncertain environment**, for **maximizing the success probability**.



Sensory Input: See a female of the same species



Agent: "Parotia Vieilloti"



Objective:
saving his genes



Azione: Dance

Environment: Forest

Learning



Sensory Input: Perceiving
the light from above

Agent: Wild Bean Plant

Action: Go up!



Feedback: Struggling

Learning



Sensory Input: Perceiving
the light from above

Agent: Wild Bean Plant

Action: Climb the pole!

Feedback: Easier, gain
more light!

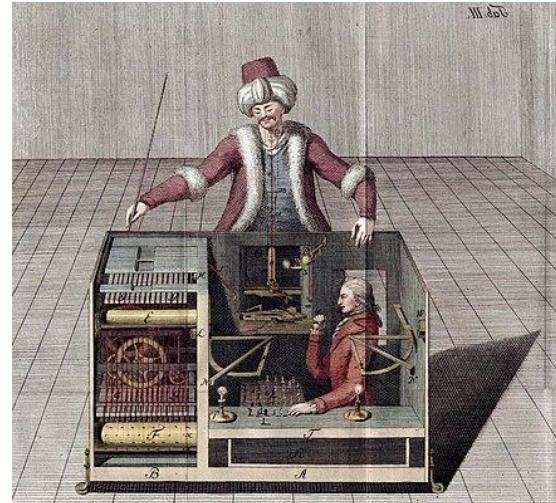
What is... “Artificial” Intelligence?



History of AI: Good Old Fashioned AI

Early Investigations on “Intelligent” Machines

- “Mimic” human behaviours



Artificial Intelligence and Computers



“The Analytical Engine has no pretensions whatever to originate anything. It can do whatever we know how to order it to perform. It can follow analysis, but it has no power of anticipating any ... relations or truths. Its province is to assist us in making available what we are already acquainted with”

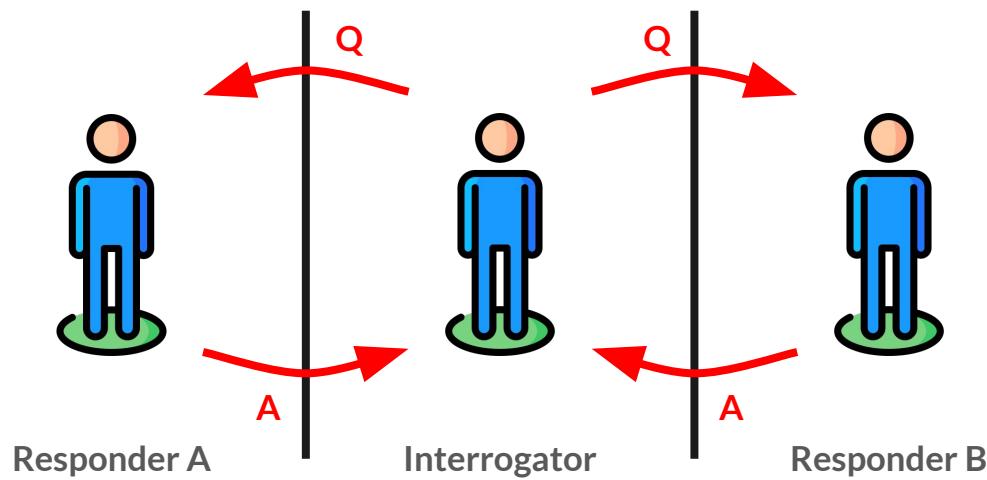
Ada Lovelace

Artificial Intelligence and Computers

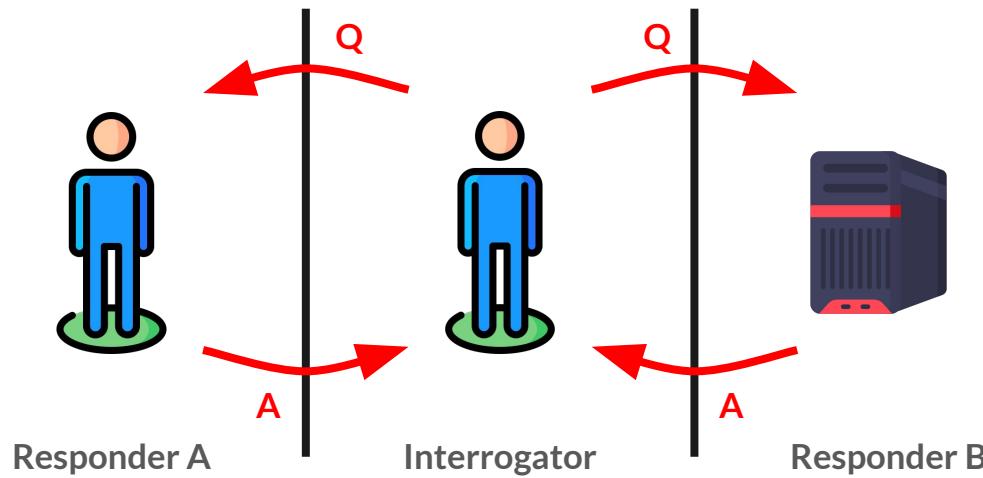
- Early Investigations on the nature of computation
- Turing Machine
- Turing Test



Turing Test




Turing Test



Who's the human and who's the machine?

Turing Test: is it reliable?

- To be intelligent does not necessarily mean to be “human”...
- But humans are intelligent



The birth of Artificial Intelligence

**1956 Dartmouth Conference:
The Founding Fathers of AI**



John MacCarthy



Marvin Minsky



Claude Shannon



Ray Solomonoff



Alan Newell



Herbert Simon



Arthur Samuel



Oliver Selfridge



Nathaniel Rochester

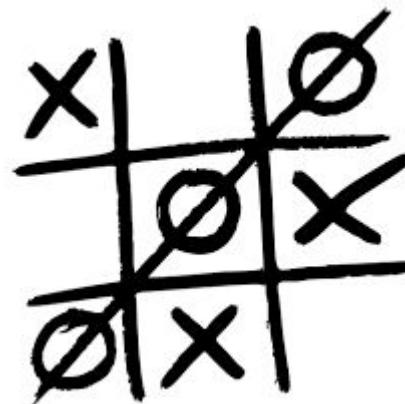


Trenchard More

AI Early Explorations: 1950s - 1960s

- Goals: identifying specific tasks that require intelligence, and figuring out how to get machines to do them
- High-Level human abilities
 - Reasoning
 - Understand Natural Language
 - Understand Images
- Low-Level abilities:
 - Recognize speech
 - Distinguishing objects in images
 - Reading cursive script

Starting point: Simple Games and Toy Problems

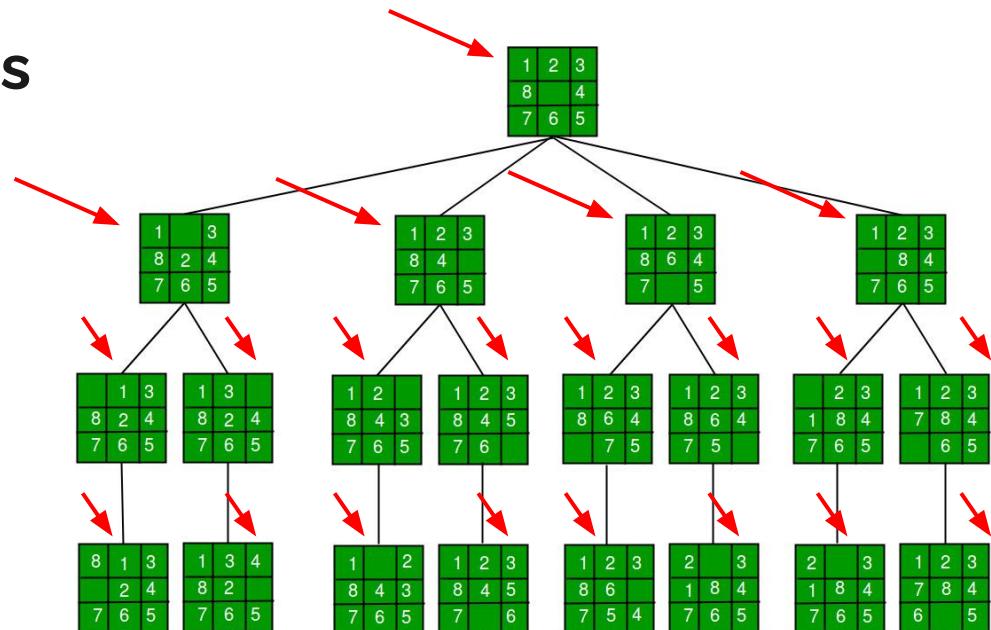


Symbols and Heuristics




Symbols and Heuristics

- Knowledge Representation: **Lists of Symbols**



Symbols and Heuristics

- Knowledge Representation: Lists of Symbols
- Search Methods: **Search tree and Heuristics**




Natural Language Processing

- Understanding, generating and translating natural language
- Difficult problem:
 - Morphology
 - Syntax
 - Semantics
 - Pragmatics
- Parse Tree

Natural Language Processing

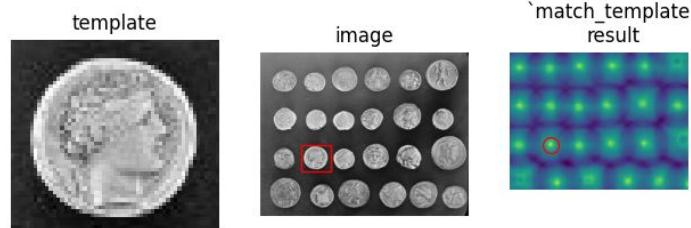
“The dog ate the bone”



- S for **sentence** (Top level)
- NP for **noun phrase**
- VP for **verb phrase** (which serves as the predicate)
- V for **verb**
- D for **determiner**
- N for **noun**

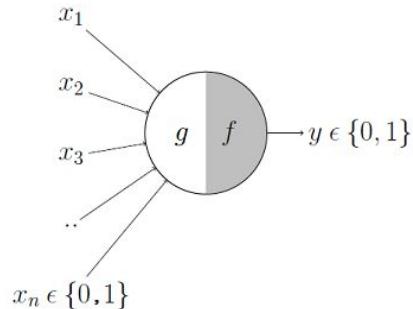
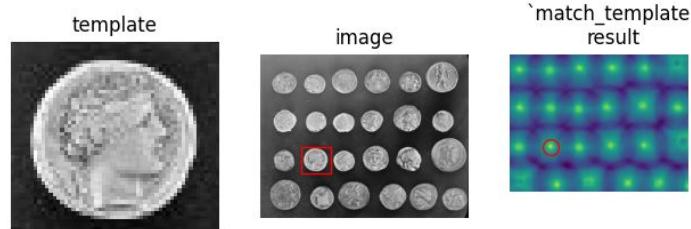
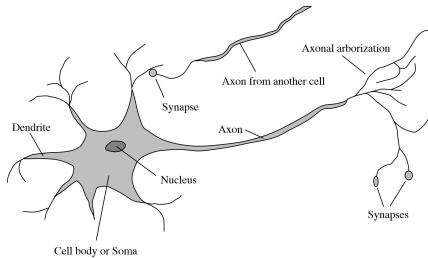
But not only Symbolic...

- Pattern Recognition
 - Template Matching
 - Feature extraction



But not only Symbolic...

- Pattern Recognition
 - Template Matching
 - Feature extraction
- McCulloch and Pitts' perceptron
 - Learning adjusting connection weights



AI Expansion: 1960s - 1980s

- From toy problems to real world commercial applications
 - Computer Vision
 - Mobile Robots
 - Game Playing
 - Speech Recognition
 - NLP
 - Knowledge Representation

AI Expansion: 1960s - 1980s - NLP

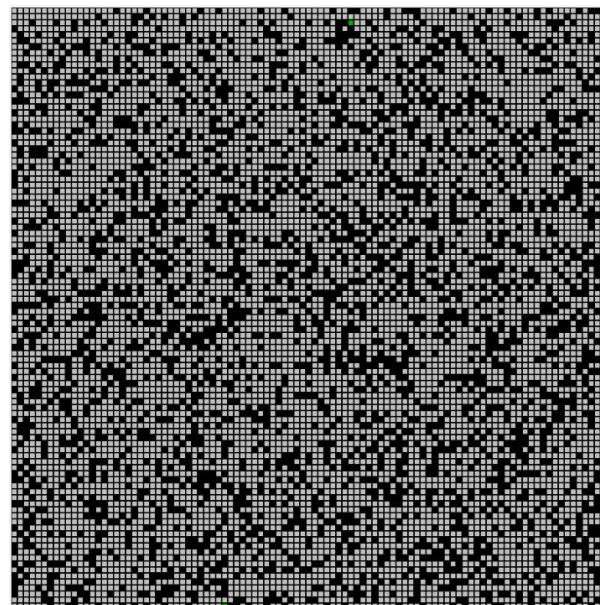
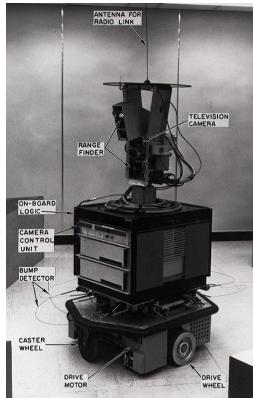
```
Welcome to
      EEEEEE  LL      IIII    ZZZZZZ  AAAAAA
      EE     LL      II      ZZ  AA  AA
      EEEE   LL      II      ZZZ  AAAAAAAA
      EE     LL      II      ZZ  AA  AA
      EEEEEE  LLLLLL IIII ZZZZZZ  AA  AA

Eliza is a mock Rogerian psychotherapist.
The original program was described by Joseph Weizenbaum in 1966.
This implementation by Norbert Landsteiner 2005.

ELIZA: Is something troubling you ?
YOU: Men are all alike.
ELIZA: What is the connection, do you suppose ?
YOU: They're always bugging us about something or other.
ELIZA: Can you think of a specific example ?
YOU: Well, my boyfriend made me come here.
ELIZA: Is it important to you that your boyfriend made you come here ?
YOU: He says I'm depressed much of the time.
ELIZA: I am sorry to hear that you are depressed.
YOU: It's true. I am unhappy.
ELIZA: Can you explain what made you unhappy ?
YOU:
```




AI Expansion: 1960s - 1980s - Mobile Robot

- Route Finding
 - A* Algorithm



AI Expansion: 1960s - 1980s

- Computer Vision



Expert Systems

- The *Good Old Fashioned AI (GOFAI)*
- Solving domain-specific problems by embedding expert knowledge in the form of logical rules
- Many success in specific tasks
- But limited achievements for “general AI”



Edward Feigenbaum

History of AI: AI Winter and the era of Big Data

AI Winter



- Real-world tasks turned out to require much more “intelligence” than that achievable by heuristic search and symbolic processing:
 - Computational Complexity
 - Background Knowledge
- Drop of interest in AI: reduced goals and less fundings



Statistical Approaches

*“Every time I fire a linguist, the performance of the speech
recognizer goes up”*

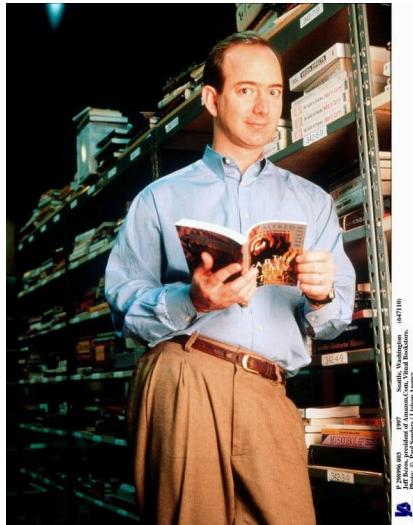
Frederick Jelinek, 1988



1980s - 1990s: from logic to statistics

- Statistical approaches seemed to overcome logical approaches
- New results in many fields
- Invented the BackPropagation Algorithm (1987)
- CIFAR10 Released (famous dataset for image classification)
- Raising Enthusiasm for statistical and inference based-AI

1990s - 2000s Data and Hardware



"Dearest Amabot,

*If you only had a heart to absorb our hatred...
Thanks for nothing, you jury-rigged rust bucket.
The gorgeous messiness of flesh and blood will prevail!"*


Seattle Weekly, 1999

1990s - 2000s Data and Hardware



- The era of data slowly begins
- Statistical methods can finally exploit large sets of data in the internet

1990s - 2000s Data and Hardware




- GPUs become progressively faster in computing data
- The combination of large data sets and powerful machines gives a boost to inference-based AI


ImageNet Dataset



Image Classification Performances



2014 - Generative Adversarial Networks



a)



b)



c)



d)

Generative Adversarial Nets

Ian J. Goodfellow,¹ Jean Pouget-Abadie,² Mehdi Mirza,¹ Bing Xu,¹ David Warde-Farley,¹ Sherjil Ozair,¹ Aaron Courville,¹ Yoshua Bengio¹
¹Département d'informatique et de recherche opérationnelle
Université de Montréal
Montréal, QC H3C 3J7

Abstract

We propose a new framework for estimating generative models via an adversarial process, in which we simultaneously train two models: a generative model G that captures the data distribution, and a discriminative model D that estimates the probability that a sample came from the training data rather than G . The training procedure for G is to maximize the probability of D making a mistake. This framework is called a generative adversarial network. We show two types of generative functions G and D : a unique solution exists, with G recovering the training data distribution and D equal to 1 everywhere. In the case where G and D are defined by multilayer perceptrons, the entire system can be trained with backpropagation. There is no need for any Markov chains or unrolled approximate inference networks during either training or generation of samples. Experiments demonstrate the potential of the framework through qualitative and quantitative evaluation of the generated samples.

1 Introduction

The promise of deep learning is to discover rich, hierarchical models [2] that represent probability distributions over the kinds of data encountered in artificial intelligence applications, such as natural images, audio waveforms containing speech, and symbols in natural language corpora. So far, the most striking successes have been in image recognition, where deep neural networks can directly map a high-dimensional, rich sensory input to a class label [14, 22]. These striking successes have primarily been based on the backpropagation and dropout algorithms, using piecewise linear units [19, 9, 10] which have a particularly well-behaved gradient. Deep generative models have had less of an impact, due to the difficulty of approximating many intractable probabilistic computations that arise in maximum likelihood estimation and related strategies, and due to difficulty of leveraging the benefits of piecewise linear units in the generative context. We propose a new generative model estimation procedure that sidesteps these difficulties.¹

In the proposed *adversarial nets* framework, the generative model is pitted against an adversary: a discriminative model that learns to determine whether a sample is from the model distribution or the data distribution. The generative model can be thought of as analogous to a team of counterfeiters, trying to produce fake currency and use it without detection, while the discriminative model is analogous to the police, trying to detect the counterfeit currency. Competition in this game drives both teams to improve their methods until the counterfeiters are indistinguishable from the genuine articles.

¹Jean Pouget-Abadie is visiting Université de Montréal from Ecole Polytechnique.

²Sherjil Ozair is visiting Université de Montréal from Indian Institute of Technology Delhi.

³Yoshua Bengio is a CIFAR Senior Fellow.

⁴All code and hyperparameters available at <http://www.github.com/goodfeli/adversarial>

2017 - Transformer Architectures (GPT)

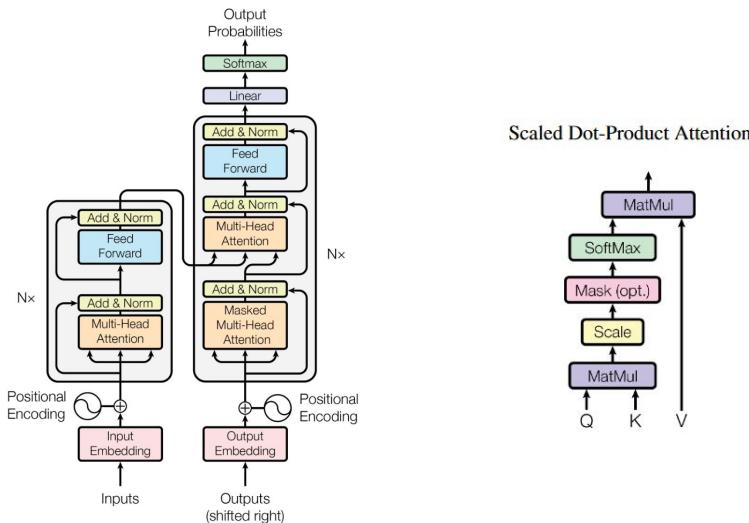

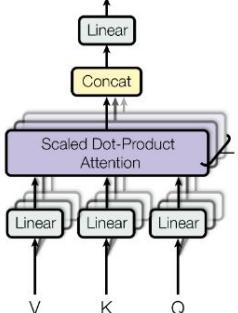


Figure 1: The Transformer - model architecture.

Scaled Dot-Product Attention



Multi-Head Attention



Provided proper attribution is provided, Google hereby grants permission to reproduce the tables and figures in this paper solely for use in journalistic or scholarly works.

Attention Is All You Need

Ashish Vaswani^{*}
Google Brain
avaswani@google.com

Noam Shazeer^{*}
Google Brain
noam@google.com

Niki Parmar^{*}
Google Research
nikip@google.com

Ilanfttuforj^{*}
Google Research
uzorf@google.com

Llion Jones^{*}
University of Toronto
llion@cs.toronto.edu

Aidan N. Gomez[†]
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser^{*}
Google Brain
lukasz.kaiser@google.com

Ilia Polosukhin[‡]
ilia-polosukhin@gmail.com

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural network architectures with many layers. These performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior, especially at longer sentence lengths, more parallelizable, and trained in less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensemble models of 28.2 BLEU. On the IWSLT 2014 English-to-German task, our model establishes a new single-model-state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best multilingual neural model. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.

^{*}Equal contribution. Listing order is random. Jakob is primarily responsible for implementing RNNs with self-attention and started this work. Ilia is primarily responsible for the multi-head attention mechanism. Noam proposed scaled dot product attention, multi-head attention and the parameter-free position representation and became the other person involved in nearly every detail of the work. Lukasz is primarily responsible for the tensorformer and tensor2tensor. Aidan and Ashish spent countless long days designing various parts of and implementing tensor2tensor, respectively, making our earlier codebase, greatly improving results and massively accelerating our research.

[†]Work performed while at Google Brain.

[‡]Work performed while at Google Research.

2021 - Zero-Shot Text-to-Image Generation

Validation



china airlines plain on the ground at an airport with baggage cars nearby.



a table that has a train model on it with other cars and things



a living room with a tv on top of a stand with a guitars sitting next to



a couple of people are sitting on a wood bench



a very cute giraffe making a funny face.



a kitchen with a fridge, stove and sink



a group of animals are standing in the snow.



Ours



Zero-Shot Text-to-Image Generation

Aditya Ramesh¹ Mikhail Pavlov¹ Gabriel Goh¹ Scott Gray¹
Chelsea Voss² Alec Radford¹ Mark Chen¹ Ilya Sutskever¹

Abstract

Text-to-image generation has traditionally focused on finding better modeling assumptions for training on a fixed dataset. These assumptions might include: automatic masking, auxiliary losses, or side information such as object part labels or segmentation masks supplied during training. We describe a simple approach for this task based on a cross-domain autoregressive model, *Text-to-Image*, and image tokens, a single token of data. With sufficient data and scale, our approach is competitive with previous domain-specific models when evaluated in a zero-shot fashion.




Figure 1: Comparison of original images, text input, and reconstruction from the discrete WAE-VAE model. The middle column shows the spatial resolution is increased by a factor of 4. White details (e.g., the fur of the cat's ear, the writing on the storefront, and the thin lines in the illustration) are sometimes lost or distorted, but the main features of the image are still typically recognizable. We use a large vocabulary size of 8192 to mitigate some of this information loss.

1. Introduction

Modern machine learning approaches to text to image synthesis started with the work of Mansimov et al. (2015), who proposed a generative adversarial network model, when extending a text encoder, that could also generate novel visual scenes. Reed et al. (2016b) later demonstrated that using a generative adversarial network (Goodfellow et al., 2014), rather than a recurrent variational autoencoder, image fidelity, Reed et al. (2016b) showed that this system could not only generate objects with recognizable properties, but also could *zero-shot* generalize to held-out categories.

Over the next few years, progress continued along a similar trajectory, including improving the generative model architecture with modifications like multi-scale generators (Zhang et al., 2017; 2018), integrating attention and auxiliary losses (Xu et al., 2018), and leveraging additional sources of conditional information beyond just text (Reed et al., 2018; 2019; 2020; Koh et al., 2017).

Separately, Nguyen et al. (2017) proposed an energy-based framework for conditional image generation that obtained a large improvement in sample quality relative to contemporary methods. Their approach can incorporate pretrained discriminative models, and they show that it is capable of performing text-to-image generation when applied to a cap-

tioning model pretrained on MS-COCO. More recently, Cho et al. (2020) also propose a method that involves optimizing the input to a pretrained cross-modal masked language model. While significant increases in visual fidelity have been claimed as a result of this work, Mansimov et al. (2015), samples can still suffer from severe artifacts such as object distortion, illogical object placement, or unnatural blending of foreground and background elements.


Recent advances fueled by large-scale generative models suggest a potential route for future improvements. Specifically, while compute and size, and data are scaling, specifically, fully autoregressive transformers (Vaswani et al., 2017) achieved impressive results in several domains such as text (Vaswani et al., 2017) and images (Chen et al., 2020), and audio (Dhariwal et al., 2020).

By comparison, text-to-image generation has typically been evaluated on relatively small datasets such as MS-COCO and CUB-200 (Welinder et al., 2010). Could dataset size and model size be the limiting factor of current approaches? In this work, we demonstrate that training a 12-billion parameter autoregressive transformer on 250 million image-text

¹OpenAI, San Francisco, California, United States. Correspondence to: Aditya Ramesh <adityar@openai.com>.

AI Today

AI Today



AI Today



EM

You

Genera una breve poesia ermetica sui culurgiones



ChatGPT

Tra le pieghe di pasta, sussurri nascosti,
Segreti di terra in bocconi avvolti.
Culurgiones, custodi di antiche storie,
In silenzio raccontano di memorie.

Sapori avvolti in mistero, sogni d'orzo e grano,
Nascosti tra i solchi di un mare lontano.
Nel loro abbraccio, il tempo si placa,
E l'anima s'immerge in dolce placida brama.

Saggi custodi di tradizioni antiche,
Racconti di terre lontane, di ricche fatiche.
In un boccone di culurgiones svelati,
Segreti dell'anima, sogni incantati.



EM

You

Traducila in sardo



ChatGPT

Sotto a sos pliegues de pasta, sussurros iscondidos,
Segretos de terra in morsi abrazados.
Culurgiones, custodes de is stòrias antichas,
In su silenzio contan de memòrias.

Sabores abrazados in misteriu, sognos de orzo e trigu,
Iscondidos entre is sulas de unu mari lontanu.
In su su so abrazu, su tempus si tranquillat,
E s'anima si ischiattat in dolce brama placida.

Saggios custodes de traditziones antichas,
Stòrias de terras lontanas, de riques fatigas.
In unu morsu de culurgiones isvelados,
Segretos de s'anima, sognos incantados.



EM

You

ora raccontala usando solo emoji



ChatGPT



AI Today



AI Today

